

Introduction to Julia as a tool in scientific computing

Luiz M. Faria

InfoMath seminar, LJLL

February 5, 2026



Table of Contents

1 Introduction to Julia

2 Case study: the N-body problem

3 Conclusions

What is Julia?

- Julia is a programming language with a **focus on scientific computing**
- First released in 2012, version 1.0 in 2018
- Open source (MIT license)
- Central promise: combine performance of C with usability of Python
- Addresses the **two-language problem**: prototype in high-level language, rewrite in low-level for performance




Using Julia

- Installation: julialang.org/downloads

```
$ curl -fsSL https://install.julialang.org | sh
```

- Basic usage: type `julia` in a terminal to open a **REPL**



- IDE: **VS Code** with `Julia` extension

 **Live demo:** `basic_usage.jl`

- Other options: Jupyter, Pluto, emacs, vim, ...

Main language features

Core features:

- Dynamic typing
- Just-In-Time (JIT) compilation via LLVM
- **Multiple dispatch** as core paradigm
- **Parametric types** for generic programming

Nice bonuses:

- Built-in package manager (`Pkg`)
- Built-in testing framework (`Test`)
- Powerful metaprogramming (macros, generated functions)
- Native Unicode support ($\alpha, \beta, \in, \dots$)
- Easy interoperability with C, Fortran, Python

Typing system

Static typing (C, Java, Rust):

- Types declared at compile time: `int x = 5;`
- Compiler checks types before running

Dynamic typing (Python, Julia):

- Types belong to *values*, not variables
- `x = 5` then `x = "hello"` is valid
- Type checked at runtime

Julia's twist

Dynamically typed, but the JIT compiler **infers types** and generates specialized code \Rightarrow flexibility + performance



Live demo: `dynamic_typing.jl`

JIT compilation

Traditional compiled (C, Fortran):

- Compile once, run many times
- Fast execution, but can slow down development cycle

Interpreted (Python, MATLAB):

- No compilation step
- Flexible, but usually slow loops

Just-In-Time (Julia):

- Compile functions **on first call**
- Specializes code based on argument types
- First call slow, subsequent calls fast



Live demo: `jit_compilation.jl`

Multiple dispatch

- Core paradigm in Julia: functions specialize on **all** argument types

Example

```
+(::Int, ::Int)    vs    +(::Float64, ::Float64)
```

Same function name, different implementations based on types

- Different from function overloading: **can** resolve at *runtime*
- Different from OOP (single dispatch)
- Most specific signature chosen
- Quite a powerful feature in practice



Live demo: `multiple_dispatch.jl`

Parametric types

- Types can have parameters:
 - `Vector{Float64}`, `Vector{Float32}`, `Vector{Int}`
 - `Matrix{T}`, `Array{T,N}`
- `Vector{T}` defines an infinite family of types parametrized by `T`
- Think template in C++
- Compiler generates specialized code for each `T` used

Key insight

Generic code + type specialization = no performance penalty



Live demo: `parametric_types.jl`

Summary so far

- We scratched the surface of Julia's core features
- To learn more I strongly suggest the [official documentation](#)
- There is also the [Julia learning page](#)
- And many other resources online...

Next: **case study** of using Julia for a scientific computing problem

Table of Contents

1 Introduction to Julia

2 Case study: the N-body problem

3 Conclusions

The N-body problem

- Classic problem in computational physics
- Simulate motion of N particles under gravitational attraction

$$m_i \ddot{\mathbf{x}}_i = - \sum_{j \neq i} \frac{G m_i m_j}{|\mathbf{x}_i - \mathbf{x}_j|^3} (\mathbf{x}_i - \mathbf{x}_j) \quad \text{for } i = 1, \dots, N$$

- Focus on the **force computation**

$$\mathbf{F}_i = - \sum_{j \neq i} \frac{G m_i m_j}{|\mathbf{x}_i - \mathbf{x}_j|^3} (\mathbf{x}_i - \mathbf{x}_j) \quad \text{for } i = 1, \dots, N$$

- Naive algorithm: double loop over all particles $\Rightarrow \mathcal{O}(N^2)$ complexity
- Better algorithms exist (Barnes-Hut, FMM) but far more complex
- Still need fast force computation kernel

Why this benchmark?

- Very simple to describe

```
function compute_forces(x, m):
    for i = 1 to N do
        for j = 1 to N, j ≠ i do
            a_i -= Gm_j  $\frac{x_i - x_j}{|x_i - x_j|^3}$ 
    return a
```

- Stress ability to efficiently handle instructions within the language
- $\mathcal{O}(N^2)$ flops with $\mathcal{O}(N)$ bytes allows for some optimizations

Why this benchmark?

- Very simple to describe

```
function compute_forces(x, m):
    for i = 1 to N do
        for j = 1 to N, j ≠ i do
            ai -= Gmj  $\frac{x_i - x_j}{|x_i - x_j|^3}$ 
    return a
```

- Stress ability to efficiently handle instructions within the language
- $\mathcal{O}(N^2)$ flops with $\mathcal{O}(N)$ bytes allows for some optimizations



A word of caution

While one may be tempted to write the algorithm above in the language of linear algebra (e.g., using matrix operations), doing so is usually *not* the most efficient approach for this problem!

Implementations

Compare three implementations of the N-body force computation:

- C

 **Live demo:** nbody.c

- Python

 **Live demo:** nbody.py

- Julia

 **Live demo:** nbody.jl

Goal is *not* to say “X is better Y”, but to show differences in performance

Inspired by <https://github.com/ahbarnett/floatingspeed/>

Performance comparison

Language	Time (s)	vs C	Gpair/s
C	0.13	1×	0.77
Julia	0.14	1.07×	0.72
Python (loops)	21	161×	0.005

Takeaway: Julia can achieve **performance comparable to C**

Next:

- How to structure larger projects?
- How to create abstractions?
- Can we make the code faster?

Using Julia (revisited)

- Scripting: write code in `.jl` files and run with `julia myscript.jl` or `include("myscript.jl")` in REPL
- Projects: put code inside a `module` and add some metadata files to help manage dependencies



Live demo: `create_pkg.jl`

Goal: move from single-file scripts to well-structured packages.

Creating a package

- Use built-in package manager `Pkg`
- Create new package skeleton with `Pkg.generate("MyPackage")`
- Main code goes in `src/MyPackage.jl`
- Tests go in `test/runtests.jl`
- Manage dependencies with `Pkg.add("DependencyName")`
- Activate project environment with `Pkg.activate(".")`

Packages for package generation

In practice people use tools like `PkgTemplates.jl` or `Bestie.jl` to automate some of the boilerplate setup (e.g. CI, documentation, tests, etc.).

Structuring the N-body code

```
1 function compute_forces()
2     ax::Vector{Float64}, ay::Vector{Float64}, az::Vector{Float64},
3     x::Vector{Float64}, y::Vector{Float64}, z::Vector{Float64},
4     m::Vector{Float64}, G::Float64 = 1.0
5     )
6     # ...
7 end
```

- ➊ Physical and logical **restructuring** of the code
- ➋ Make types **parametric** for flexibility
- ➌ Optimize for **performance**

```
1 struct GravitationalSystem
2     positions::Matrix{Float64}
3     acceleration::Matrix{Float64}
4     masses::Vector{Float64}
5     G::Float64
6 end
7
8 function compute_forces!(sys::GravitationalSystem)
9     # Compute gravitational forces and update accelerations
10    # ...
11 end
```



```
1 struct GravitationalSystem{T <: Real}
2     positions::Matrix{T}
3     acceleration::Matrix{T}
4     masses::Vector{T}
5     G::T
6 end
7
8 function compute_forces!(sys::GravitationalSystem)
9     # Compute gravitational forces and update accelerations
10    # ...
11 end
```



```
1 struct GravitationalSystem{T <: Real}
2     positions::Matrix{T}
3     acceleration::Matrix{T}
4     masses::Vector{T}
5     G::T
6 end
7
8 function compute_forces!(sys::GravitationalSystem{T}) where {T}
9     N = pick_vector_width(T) # register width
10    # Compute gravitational forces and update accelerations
11    # ...
12 end
```

1. Organize

2. Generalize

3. Optimize



Live demo: `gravitational_system.jl`

Going faster

The simple implementation is not bad, but we can do better

- ✓ Memory layout and allocations
- ✗ Use of vectorized instructions (SIMD)
- ✗ Parallelization (multi-threading)

SIMD instructions

Single Instruction, Multiple Data

Scalar: $a_1 + b_1 = c_1$

SIMD: $\begin{matrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{matrix} + \begin{matrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{matrix} = \begin{matrix} c_1 \\ c_2 \\ c_3 \\ c_4 \end{matrix}$

In Julia:

- `@simd hints`
- `@turbo (LoopVectorization.jl)`
- Explicit SIMD via `SIMD.jl`

Hardware (processor architectures)


- x86-64: SSE, AVX, AVX2, AVX-512, ...
- ARM: NEON, SVE, ...
- ...



Live demo: `simd.jl`

Multi-threading

Execute work in parallel across multiple CPU cores



In Julia:

- Start with `julia -t auto` or `julia -t 8`
- `Threads.@threads` for loop parallelism
- `@spawn` / `@sync` for task-based parallelism

Considerations

- Race conditions: use locks or atomic operations
- Overhead: threading has a cost, use for large workloads
- N-body force computation is **embarrassingly parallel**

Optimizations

Next:

- Try to **vectorize** N-body force computation with SIMD
- Add **multi-threading** to parallelize over particles



Live demo: gravitational_system.jl

Version	Gpairs/s	Speedup
Julia (plain)	0.65	1×
Julia (SIMD)	1.5	2.3
Julia (SIMD + threads)	8.58	12.4

Hardware

- Vector register width: 128-bit
- Number of threads used: 6

Bonus: reduced precision

- Our code is already generic over floating point types
- So does it work with e.g. `Float32`? `Float16`? Faster?



Live demo: `bench.jl`

Version	Gpairs/s	Speedup
Julia (SIMD + threads + <code>Float64</code>)	8.58	1×
Julia (SIMD + threads + <code>Float32</code>)	17.62	2.05
Julia (SIMD + threads + <code>Float16</code>)	34.89	4.06

Bonus: reduced precision

- Our code is already generic over floating point types
- So does it work with e.g. `Float32`? `Float16`? Faster?



Live demo: `bench.jl`

Version	Gpairs/s	Speedup
Julia (SIMD + threads + <code>Float64</code>)	8.58	1×
Julia (SIMD + threads + <code>Float32</code>)	17.62	2.05
Julia (SIMD + threads + <code>Float16</code>)	34.89	4.06



A word of caution

- Effect of round-off seeing much earlier in e.g. `Float32`, `Float16`
- Stability considerations may require algorithmic changes
- Mixed-precision algorithms are an active research area

Table of Contents

1 Introduction to Julia

2 Case study: the N-body problem

3 Conclusions

Summary

Explored how Julia can be used for **number-crunching** tasks

. Takeaways:

- ✓ Possible to write fast code in Julia without leaving the language
- ✓ Key features: JIT compilation, multiple dispatch, parametric types
- ✓ Implementation matters: data layout, SIMD, parallelism, ...
- ✓ Easier to learn and use than many traditional compiled languages
- ✓ An interesting language choice for scientific computing projects

Where Julia falls short

I talked about some strengths of Julia, but it's **not perfect**. Some pain points to be aware of:

- ✗ Latency of JIT compilation can be annoying for scripting workflow
- ✗ External tooling ecosystem still maturing (debuggers, profilers, ...)
- ✗ Type instabilities are a common source of performance issues
- ✗ Lack of strict interface makes it hard to enforce API contracts
- ✗ Hard to compile to a standalone binary (improving!)

Suggestion: try it out on a small project and draw your own conclusions!



<https://julialang.org/>

Questions?